

Problem 1: Can you fit a line?

Problem Statement

Consider N data points (\mathbf{x}_i, y_i) for $i = 1, \dots, N$ obtained from an experiment, where $\mathbf{x}_i \in \mathbb{R}^n$ and $y_i \in \mathbb{R}$. Let $N > n + 1$. The goal of linear regression is to find the “best” linear fit; i.e. find $\mathbf{c} \in \mathbb{R}^n$, $d \in \mathbb{R}$ s.t.:

$$y_i \approx \mathbf{c}^T \mathbf{x}_i + d, \text{ for } i = 1, \dots, N$$

a

Suppose each measurement is corrupted by independent Gaussian noise with identical variances and 0 mean. Find $\mathbf{A} \in \mathbb{R}^{N \times (n+1)}$ in terms of $\mathbf{x}_1, \dots, \mathbf{x}_N$ s.t. that the maximum likelihood estimate of $\hat{\mathbf{c}}, \hat{d}$ is

$$\begin{pmatrix} \hat{\mathbf{c}} \\ \hat{d} \end{pmatrix} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}, \text{ where } \mathbf{y} := \begin{pmatrix} y_1 \\ \vdots \\ y_N \end{pmatrix}$$

b

Consider the case with $n = 1$; i.e. x_1, \dots, x_N are scalars. Finding the best linear fit amounts to finding (\hat{c}, \hat{d}) that minimizes:

$$J(c, d) := \sum_{i=1}^N (y_i - cx_i - d)^2$$

Compute a stationary point (\hat{c}, \hat{d}) by setting the derivative of J w.r.t c and d to zero.

c

Use the second derivative tests to conclude that your (\hat{c}, \hat{d}) computed in (b) is indeed a local minimizer of J . Can you conclude from this second derivative test alone that it is a global minimizer?

d

Consider the following (x, y) pairs:

- (1.00, 1.10)
- (1.50, 1.62)
- (2.00, 1.98)
- (2.57, 2.37)
- (3.00, 3.23)
- (3.50, 3.69)
- (4.00, 3.97)

Draw a scatter plot of these points; then, plot the best linear fit to these points using your formula in part (b).

Solution**a**

We can rewrite the given equation to:

$$y_i = \mathbf{x}_i^T \mathbf{c} + d$$

Additionally, let: $\mathbf{1}_N$ be a column vector of 1's with N rows. Thus:

$$\mathbf{y} = \begin{pmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_N \end{pmatrix} \mathbf{c} + d \mathbf{1}_N$$

We can further re-arrange this to:

$$\mathbf{y} = \begin{pmatrix} \mathbf{x}_1 & 1 \\ \mathbf{x}_2 & 1 \\ \vdots & \vdots \\ \mathbf{x}_N & 1 \end{pmatrix} \begin{pmatrix} \mathbf{c} \\ d \end{pmatrix} \Rightarrow \mathbf{A} = \begin{pmatrix} \mathbf{x}_1 & 1 \\ \mathbf{x}_2 & 1 \\ \vdots & \vdots \\ \mathbf{x}_N & 1 \end{pmatrix}$$

b

$$\nabla J(c, d) = \begin{pmatrix} \frac{\partial J}{\partial c} \\ \frac{\partial J}{\partial d} \end{pmatrix} = 2 \begin{pmatrix} \sum_{i=1}^N (y_i - cx_i - d)(-x_i) \\ \sum_{i=1}^N (y_i - cx_i - d)(-1) \end{pmatrix}$$

Setting equal to zero and noting that the 2 can then just drop out (and push the negative sign around):

$$\mathbf{0} = \begin{pmatrix} \sum_{i=1}^N (cx_i + d - y_i)(x_i) \\ \sum_{i=1}^N (cx_i + d - y_i) \end{pmatrix}$$

Expanding the top:

$$0 = \sum_{i=1}^N (cx_i^2 + dx_i - y_i x_i) = c \left(\sum x_i^2 \right) + d \left(\sum x_i \right) - \left(\sum x_i y_i \right)$$

And the bottom:

$$0 = \sum_{i=1}^N (cx_i + d - y_i) = c \left(\sum x_i \right) + d \left(\sum 1 \right) - \left(\sum y_i \right)$$

Solve for c :

$$c = \left(\left(\sum y_i \right) - d \left(\sum 1 \right) \right) / \left(\sum x_i \right)$$

Note that $\sum_{i=1}^N 1 = N$. Plug into the first partial derivative equation:

$$0 = \frac{\sum y_i - Nd}{\sum x_i} \sum x_i^2 + d \sum x_i - \sum x_i y_i$$

Solve for d :

$$\frac{\sum x_i^2 \sum y_i - Nd \sum x_i^2}{\sum x_i} + d \sum x_i - \sum x_i y_i$$

$$\begin{aligned}
&= \sum x_i^2 \sum y_i - Nd \sum x_i^2 + d \left(\sum x_i \right)^2 - \sum x_i \sum x_i y_i \\
Nd \sum x_i^2 - d \left(\sum x_i \right)^2 &= \sum x_i^2 \sum y_i - \sum x_i \sum x_i y_i \\
d \left(N \sum x_i^2 - \left(\sum x_i \right)^2 \right) &= \sum x_i^2 \sum y_i - \sum x_i \sum x_i y_i \\
d &= \frac{\sum x_i^2 \sum y_i - \sum x_i \sum x_i y_i}{N \sum x_i^2 - \left(\sum x_i \right)^2}
\end{aligned}$$

which can then be plugged back in to get c . I'm lazy so I used Wolfram to simplify the expression and got:

$$c = \frac{\sum x_i \sum y_i - N \sum x_i y_i}{\left(\sum x_i \right)^2 - N \sum x_i^2}$$

c

Given the gradient, the Hessian is trivially calculable as:

$$\nabla^2 J(c, d) = 2 \begin{pmatrix} \sum x_i^2 & \sum x_i \\ \sum x_i & N \end{pmatrix}$$

If $\nabla^2 J \succcurlyeq 0$ everywhere then J is convex and the first order condition ($\nabla J = 0$) is sufficient for both a global and local optimality. HOWEVER, this will ultimately depend on the x 's chosen, as well as the quantity N chosen. The eigenvalues are given by¹:

$$\lambda_{1,2} = \frac{1}{2} \left(\sum x_i^2 + N \pm \sqrt{\left(\sum x_i^2 \right)^2 + 4 \sum (x_i)^2 + N^2 - 2N \sum x_i^2} \right)$$

which must be $\lambda_{1,2} \geq 0$ for the matrix to be PSD. We know that $N \geq 1$ and $x_i \in \mathbb{R} \implies \sum x_i^2 \geq 0$, so for both eigenvalues to be non-negative the following must hold:

$$\sum x_i^2 + N \geq \sqrt{\sum x_i^2 - 2N \sum x_i^2 + 4 \left(\sum x_i \right)^2 + N^2}$$

Failing this condition, the matrix is not PSD and we cannot assume global optimality (although it doesn't imply it is NOT a global optimum).

¹Thanks, Wolfram!

d

Code is made available at the end of document.




Figure 1: Linear Regression on the provided data, using both the derived expression (dashed orange line) and SciPy's built-in method for verification (solid blue line)

Problem 2: Estimate with Confidence

Problem Statement

Suppose the true parameters of a system are described by $x = (1, 1)^T$, and the measurements are given by:

$$\mathbf{y} = \underbrace{\begin{pmatrix} 1 & 2 \\ 2 & 3 \\ 3 & 4 \\ 4 & 5 \end{pmatrix}}_{:=\mathbf{A}} \mathbf{x} + \mathbf{e}$$

where $e \in \mathbb{R}^4$ is sampled from a multivariate Gaussian distribution with zero mean and covariance matrix $\mathbf{R} := \text{diag}(0.1, 0.2, 0.3, 0.4)$. That is, $\mathbf{e} \sim \mathcal{N}(0, \mathbf{R})$.

a

Write code for the following experiment:

1. Generate an error vector \mathbf{e} according to $\mathcal{N}(0, \mathbf{R})$. Compute \mathbf{y} from it and report it.
2. In class we derived a formula to compute the maximum likelihood estimate $\hat{\mathbf{x}}$ from \mathbf{y} . Using \mathbf{y} from the previous step, compute and report your estimate $\hat{\mathbf{x}}$.
3. The estimated error is given by $\mathbf{y} - \mathbf{A}\hat{\mathbf{x}}$. With \mathbf{y} from the prior steps, compute and report:

$$J(\hat{\mathbf{x}}) := (\mathbf{y} - \mathbf{A}\hat{\mathbf{x}})^T \mathbf{R}^{-1} (\mathbf{y} - \mathbf{A}\hat{\mathbf{x}})$$

Recall that we claimed (without proof) that $J(\hat{\mathbf{x}}) \sim \chi_2^2$. Then, we have:

$$\mathbb{P}\{J(\hat{\mathbf{x}}) \leq 5.9915\} = 0.95$$

where $\mathbb{P}\{\mathcal{E}\}$ denotes the probability of an event \mathcal{E} . Conclude with a 95% confidence level whether your estimate error conforms to the error model we assumed.

b

Repeat the experiment in the prior section TEN-THOUSAND TIMES. Each time you run the experiment, record the value of $J(\hat{\mathbf{x}})$. Plot a histogram of $J(\hat{\mathbf{x}})$. Comment qualitatively how the histogram of $J(\hat{\mathbf{x}})$ compares to the PDF of a χ_2^2 random variable. Also, report the percentage of times you concluded that your measurement did not conform to the error model. Compare this fraction to your confidence interval in the prior portion.

Solution

a

Again, code is made available at the end of the assignment. The output of the script is:

```
First y:
[2.93953942  4.85114251  7.50790296  8.75054375]
x hat:
[1.1427999  0.90150926]
J: 1.1135043299064333, Conforms to error model (with 95% confidence): True
```

b

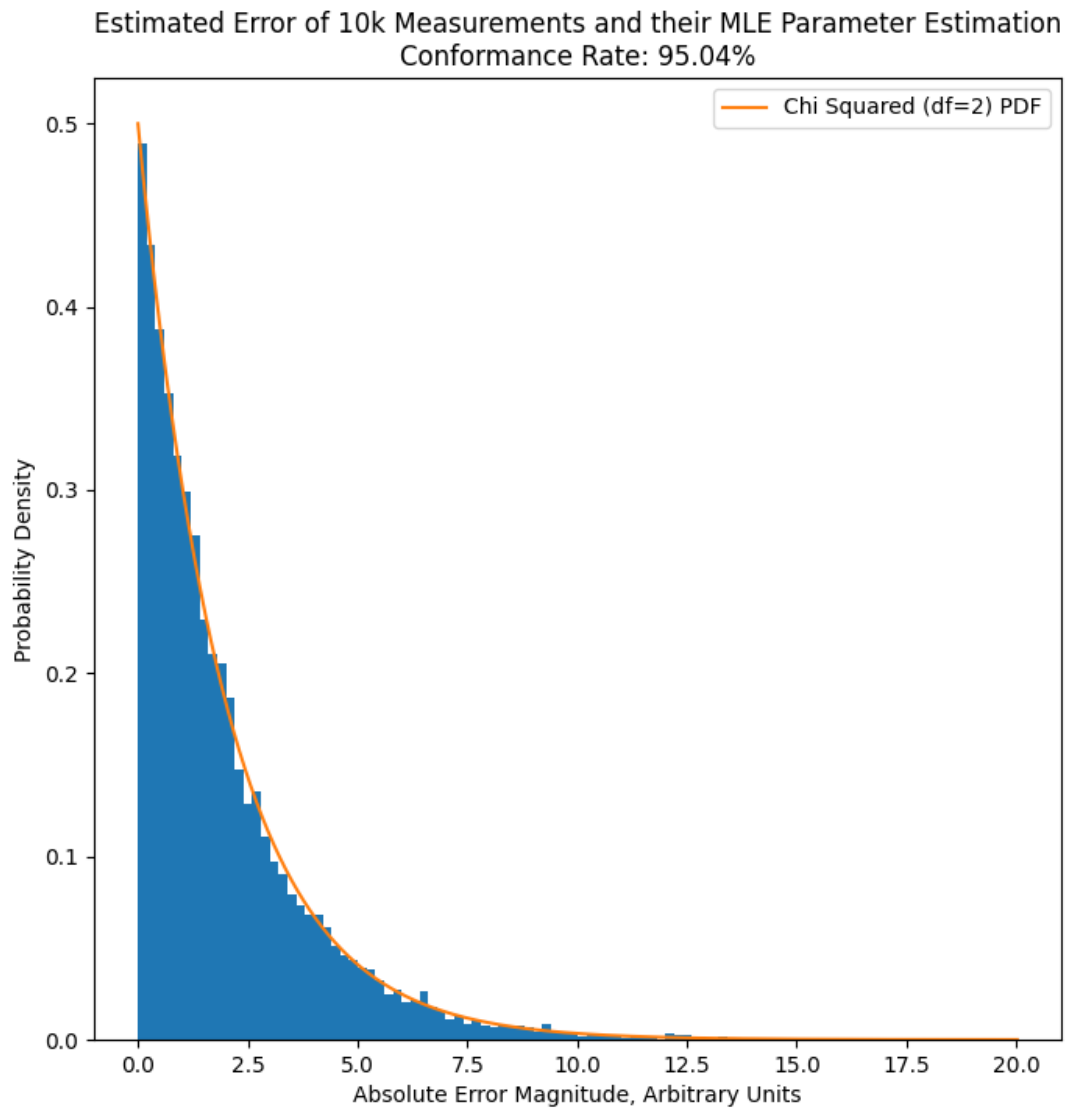


Figure 2: Density based histogram of 10k experiments with a 2-DoF Chi-Squared PDF superimposed. Qualitatively...pretty dang good!

Code

```

1 from typing import Callable
2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from scipy.stats import linregress, chi2
6
7 DATA = np.array(
8     [
9         (1.00, 1.10),
10        (1.50, 1.62),
11        (2.00, 1.98),
12        (2.50, 2.37),
13        (3.00, 3.23),
14        (3.50, 3.69),
15        (4.00, 3.97),
16    ]
17 ).T
18
19
20 def scipy_regression(
21     x: np.ndarray, y: np.ndarray
22 ) -> Callable[[np.ndarray], np.ndarray]:
23     r = linregress(x, y)
24     f = lambda x: r.slope * x + r.intercept
25
26     return f
27
28
29 def my_regression(x: np.ndarray, y: np.ndarray) -> Callable[[np.ndarray], np.ndarray]:
30     N = y.size
31     d = (np.sum(x**2) * np.sum(y) - np.sum(x) * np.sum(x * y)) / (
32         N * np.sum(x**2) - np.sum(x) ** 2
33     )
34     c = (np.sum(x) * np.sum(y) - N * np.sum(x * y)) / (
35         np.sum(x) ** 2 - N * np.sum(x**2)
36     )
37
38     return lambda x: c * x + d
39
40
41 def problem_one() -> None:
42     x = DATA[0]
43     y = DATA[1]
44     f_hat = scipy_regression(x, y)
45     f_mine = my_regression(x, y)
46     x_base = np.linspace(0, 5, 100)
47     y_hat_scipy = f_hat(x_base)
48     y_hat_mine = f_mine(x_base)
49     plt.figure(figsize=(8, 8))
50     plt.title("Linear Regressions")
51     plt.scatter(x, y, label="Raw Data", c="red")
52     plt.plot(x_base, y_hat_scipy, label="SciPy Regression")
53     plt.plot(
54         x_base, y_hat_mine, label="Explicitly calculated Regression", linestyle="dashed"
55     )
56     plt.xlabel("x")
57     plt.ylabel("y")
58     plt.legend()
59
60     plt.savefig("./hw3_regression.png")
61
62
63 def generate_e(mean: np.ndarray, covariance_matrix: np.ndarray) -> np.ndarray:
64     return np.random.multivariate_normal(mean, cov=covariance_matrix)
65
66
67 class NoisySystem:
68     def __init__(
69         self, A: np.ndarray, x: np.ndarray, mean: np.ndarray, R: np.ndarray
70     ) -> None:
71         self.A = A
72         self.x = x
73         self.mean = mean
74         self.R = R
75
76     def __call__(self) -> np.ndarray:
77         return self.A @ self.x + generate_e(self.mean, self.R)
78
79     def __repr__(self) -> str:
80         return f"A:\n{self.A}\nx:\n{self.x}\nMean:\n{self.mean}\nR:\n{self.R}"
81
82
83 class LinearEstimator:
84     def __init__(self, A: np.ndarray, R: np.ndarray) -> None:
85         # I know, I know, directly calling an inverse is bad...
86         self.R_inv = np.linalg.inv(R)
87         self.A = A
88         self.da_matrix = np.linalg.inv(A.T @ self.R_inv @ A) @ A.T @ self.R_inv
89

```

```

90     def __call__(self, y: np.ndarray) -> np.ndarray:
91         return self.da_matrix @ y
92
93     def estimate_error(self, y: np.ndarray, x_hat: np.ndarray) -> float:
94         tmp = y - (self.A @ x_hat)
95         return tmp.T @ self.R_inv @ tmp
96
97
98 def problem_2_experiment(system: NoisySystem, estimator: LinearEstimator) -> float:
99     y = system()
100     x_hat = estimator(y)
101     return estimator.estimate_error(y, x_hat)
102
103
104 def problem_two() -> None:
105     mean = np.zeros(4)
106     R = np.diag([0.1, 0.2, 0.3, 0.4])
107
108     A = np.array([[1, 2], [2, 3], [3, 4], [4, 5]])
109     x_true = np.array([1, 1])
110     system = NoisySystem(A, x_true, mean, R)
111     estimator = LinearEstimator(A, R)
112
113     first_y = system()
114     print(f"First y:\n{first_y}")
115     x_hat = estimator(first_y)
116     print(f"x_hat:\n{x_hat}")
117     J = estimator.estimate_error(first_y, x_hat)
118     print(f"J: {J}, Conforms to error model (with 95% confidence): {J <= 5.9915}")
119
120     EXPERIMENT_COUNT = 10000
121     rslts = np.zeros(EXPERIMENT_COUNT)
122     for i in range(EXPERIMENT_COUNT):
123         rslts[i] = problem_2_experiment(system, estimator)
124
125     times_conformed = np.count_nonzero(rslts <= 5.9915)
126     error_rate = (times_conformed / EXPERIMENT_COUNT) * 100
127
128     df = 2
129     x = np.linspace(rslts.min(), rslts.max(), 10000)
130
131     fig, ax = plt.subplots(1, 1, figsize=(8, 8))
132     ax.hist(rslts, bins=100, density=True)
133     ax.set_title(
134         f"Estimated Error of 10k Measurements and their MLE Parameter Estimation\nConformance Rate: {
135             error_rate}%")
136     ax.set_xlabel("Absolute Error Magnitude, Arbitrary Units")
137     ax.set_ylabel("Probability Density")
138     ax.plot(x, chi2(df).pdf(x), label="Chi Squared (df=2) PDF")
139     ax.legend()
140     plt.savefig("./problem2_hist.png")
141
142
143 if __name__ == "__main__":
144     problem_one()
145     problem_two()

```