



Problem 1: Appreciate Cholesky Graphically

Problem Statement



(a) Graph G_1



(b) Graph G_2

a

Find a matrix $\mathbf{A} \in \mathbb{R}^{5 \times 5}$ s.t. $\mathcal{G}(\mathbf{A}) = G_1$ in the figure above, and s.t. $\mathbf{A} \succ 0$. *Hint: Associate positive numbers or weights to each node and each edge of G_1 . Define $A_{ij} = A_{ji}$ as the negative of the weight on the edge (i, j) of G_1 . Define A_{ii} to be the sum of the weight of node i and the weights of all edges connected to node i . Such a matrix is called a Laplacian matrix for G_1 . Compute the eigenvalues to verify it is indeed PD.*

b

Compute the lower triangular Cholesky factor $\mathbf{L} \in \mathbb{R}^{5 \times 5}$ with positive diagonals such that $\mathbf{L}\mathbf{L}^T$ equals \mathbf{A} in part (a). Draw a graph on 5 nodes that describes the sparsity pattern of \mathbf{L} ; i.e. draw $G(\mathbf{L})$. Verify that $G(\mathbf{L})$ is a chordal graph.

c

Notice that G_2 in figure 1b is a chordal graph. Find a permutation matrix \mathbf{P} s.t. that when you compute the Cholesky factor \mathbf{L}' of $\mathbf{P}\mathbf{A}\mathbf{P}^T$ with your \mathbf{A} from part (a), then $G(\mathbf{L}') = G_2$. *Hint: Eliminate the nodes of G_1 in the sequence $(5, 2, 3, 1, 4)$. Encode that elimination order in a permutation matrix.*

d()**

Let $\mathbf{X} \in \mathbb{R}^{n \times n}$ be a PD matrix and $\mathbf{\Gamma} \in \mathbb{R}^{n \times n}$ be a lower triangular matrix with positive diagonals, s.t. $\mathbf{X} = \mathbf{\Gamma}\mathbf{\Gamma}^T$. Prove that $G(\mathbf{X})$ is a subgraph of $G(\mathbf{\Gamma})$ and $G(\mathbf{\Gamma})$ is a chordal graph. *Hint: Use an induction argument on the size of the matrix. Characterize how the steps in Cholesky factorization shapes the sparsity pattern of $\mathbf{\Gamma}$. Finally, utilize the fact that if a node and a set of edges from that node is added to a chordal graph G , then the obtained graph is chordal IFF every two neighbors of the new node in G already shared an edge in G .*

Solution

For script, see relevant section at the end of the document.

a

Relevant script output:

```
1a Laplacian:
[[ 3. -1. -1.  0.  0.]
 [-1.  3.  0. -1.  0.]
 [-1.  0.  4. -1. -1.]
 [ 0. -1. -1.  3.  0.]
 [ 0.  0. -1.  0.  2.]]
1a eigenvalues:
[5.481  1.    1.83   3.    3.689]
1a PD: True
```




Figure 1: Graph G_1 generated in networkx, for verification

b

```
1b L:
[[ 1.732  0.    0.    0.    0.   ]
 [-0.577  1.633  0.    0.    0.   ]
 [-0.577 -0.204  1.904  0.    0.   ]
 [ 0.    -0.612 -0.591  1.509  0.   ]
 [ 0.     0.   -0.525 -0.206  1.297]]
```




Figure 2: Graph $\mathcal{G}(L)$ generated in networkx. Note it is indeed chordal.


c

Figure 3: Graph $\mathcal{G}(L')$ generated in networkx. Note it indeed matches the structure of G_2 with the nodes swapped as described.

d

...meh, it's 4PM on the Friday before break.

Problem 2: Split 'em matrices for Linear Systems

Problem Statement

Consider a linear system of equations $Ax = b$ where $A \in \mathbb{R}^{n \times n}$ and $\exists A^{-1}$. The matrix splitting method splits A as $A = M - N$ and successively solves the equation:

$$x^{k+1} = M^{-1} (Nx^k + b) \quad (1)$$

for $k \geq 0$ to perform a fixed-point iteration in order to solve the linear system $Ax = b$. Identify L, D and U as the strictly lower triangular, diagonal, and strictly upper triangular part of A . Then, $A = L + D + U$. When $M = L + D$, then the fixed point iteration becomes:

$$x_i^{k+1} := \frac{1}{A_{ii}} \left[b_i - \sum_{j < i} A_{ij} x_j^{k+1} - \sum_{j > i} A_{ij} x_j^k \right] \quad (2)$$

for $i = 1, \dots, n$ and $k \geq 0$.

a

Write the update equation for each element of x as in equation 2 when $M = D + U$ using backward induction to compute the iterates in equation 1 successively.

b

Call the method you devised in part (a) as the *backward* Gauss-Seidel method. Which among this method and Jacobi methods would you expect to converge faster to the solution of $Ax = b$? Which one will you choose if distributed computation is involved, and why?

c

Implement both backward Gauss-Seidel and Jacobi methods to solve the linear system of equations $Ax = b$ with

$$A := \begin{pmatrix} 10 & 5 & 3 & 4 \\ 4 & 10 & 2 & 1 \\ 1 & 3 & 8 & 2 \\ 1 & 6 & 3 & 9 \end{pmatrix}, \quad b := \begin{pmatrix} 4 \\ -5 \\ 4 \\ -11 \end{pmatrix}$$

starting with $x^0 := (1, 1, 1, 1)^T$. Plot the residues and errors in the successive iterates on a semilog plot. Iterate until the residue falls below 10^{-5} . The error and residue, respectively, at iteration k are given by:

$$\epsilon^k := \|x - x^*\|, \quad r^k := \|b - Ax^k\|$$

where x^* is the solution to the linear system.

d

Recall that the method of successive over-relaxation (SOR) uses $M = \frac{1}{\omega} D + L$. Any matrix splitting method converges to the solution of a linear system *IF* the spectral radius of the matrix $I - M^{-1}A$ is less than 1. Plot the spectral radius of $I - M^{-1}A$ for SOR as a function of ω in the range $[0.01, 2.20]$ for 20 randomly generated PD A matrices on the same graph. Based on your plot, can you guess for what values of ω SOR converges for PD matrices?

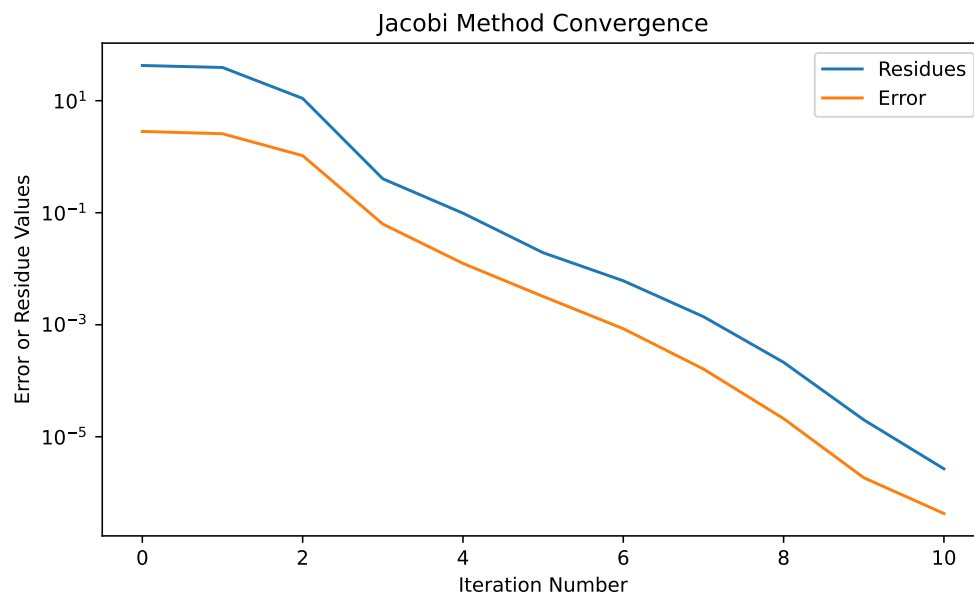
Solution**a**

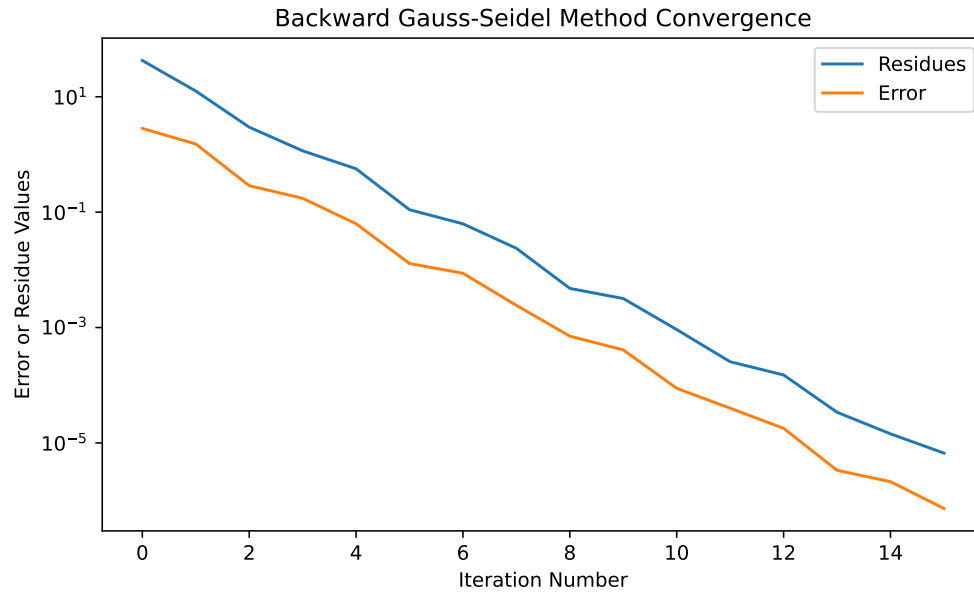
$$x_i^{k+1} := \frac{1}{A_{ii}} \left[b_i - \sum_{j>i} A_{ij} x_j^{k+1} - \sum_{j<i} A_{ij} x_j^k \right]$$

b


It will depend. Gauss-Seidel methods use more “up to date” information, and so intuitively we would expect faster convergence. However, this is not a hard and fast rule, and there are counterexamples wherein one method converges and the other doesn’t, OR that one will converge faster in certain cases, or slower in others. See [1].

I think that using a typical scatter-gather scheme for distributed computation, Jacobi iteration would be preferred for distributed computing. Each element of x^{k+1} can be calculated in an embarrassingly parallel fashion (i.e. each element can be calculated in a fully independent fashion), whereas Gauss-Seidel will require iterations and sharing current estimates of x^{k+1} around, which will limit the ability to parallelize.

c



d

Spectral Radius of the iteration matrix $I-M^{-1}A$ for 20 random A matrices

Empirically it appears that the spectral radius will be less than 1 for values between 0 and 2. This is in fact a proven property of SOR methods, according to the Wikipedia article on them.

Code

Problem 1

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 import networkx as nx
4
5 label_conversion_dict = {0: 1, 1: 2, 2: 3, 3: 4, 4: 5}
6
7 np.set_printoptions(precision=3)
8
9
10 def get_G1() -> nx.Graph:
11     g = nx.Graph()
12     g.add_nodes_from((1, 2, 3, 4, 5))
13     g.add_edges_from(((1, 2), (2, 4), (4, 3), (3, 5), (3, 1)))
14
15     return g
16
17
18 def get_G2() -> nx.Graph:
19     g = get_G1()
20     g.add_edge(1, 4)
21
22     return g
23
24
25 def plot_graph(G, label_conversion=label_conversion_dict):
26     pos = nx.planar_layout(G)
27     nx.draw_networkx_nodes(G, pos)
28     nx.draw_networkx_edges(G, pos)
29     nx.draw_networkx_labels(G, pos, label_conversion)
30
31
32 def one_a():
33     # The definition of the Laplacian given by the HW assigns a weight to the
34     # node as well, so we must add some positive diagonal matrix (use I for
35     # convenience)
36     laplacian = nx.laplacian_matrix(get_G1()).toarray() + np.eye(5)
37     print(f"1a Laplacian:\n{laplacian}")
38     eigvals = np.linalg.eigvals(laplacian)
39     print(f"1a eigenvalues:\n{eigvals}")
40     print(f"1a PD: {not np.any(np.isclose(eigvals, 0.0)) and np.all(eigvals > 0)}")
41     plot_graph(get_G1(), label_conversion=None)
42     plt.savefig("one.a.pdf")
43
44
45 def one_b():
46     laplacian = nx.laplacian_matrix(get_G1()).toarray() + np.eye(5)
47     L = np.linalg.cholesky(laplacian)
48     print("1b L:")
49     print(L)
50     G_of_L = nx.from_numpy_array(L, create_using=nx.Graph)
51     G_of_L.remove_edges_from(nx.selfloop_edges(G_of_L))
52     plt.figure()
53     plot_graph(G_of_L)
54     plt.savefig("one.b.pdf")
55
56
57 def one_c():
58     P = np.array(
59         (
60             (0, 0, 0, 0, 1),
61             (0, 1, 0, 0, 0),
62             (0, 0, 1, 0, 0),
63             (1, 0, 0, 0, 0),
64             (0, 0, 0, 1, 0),
65         )
66     )
67     laplacian = nx.laplacian_matrix(get_G1()).toarray() + np.eye(5)
68     reordered = P @ laplacian @ P.T
69     L = np.linalg.cholesky(reordered)
70     G_of_L = nx.from_numpy_array(L, create_using=nx.Graph)
71     G_of_L.remove_edges_from(nx.selfloop_edges(G_of_L))
72     plt.figure()
73     plot_graph(G_of_L, label_conversion=label_conversion_dict)
74     plt.savefig("one.c.pdf")
75
76
77 if __name__ == "__main__":
78     one_a()
79     one_b()
80     one_c()

```


Problem 2

```

1 from typing import Tuple, List
2
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 A = np.array(
7     (
8         (10, 5, 3, 4),
9         (4, 10, 2, 1),
10        (1, 3, 8, 2),
11        (1, 6, 3, 9),
12    )
13 )
14
15 b = np.array((4, -5, 4, -11))
16 x_exact = np.linalg.solve(A, b)
17 print(x_exact)
18 err_func = lambda xk: np.linalg.norm(xk - x_exact)
19
20
21 def residual(A: np.ndarray, x: np.ndarray, b: np.ndarray) -> float:
22     return np.linalg.norm(b - A @ x)
23
24
25 def Jacobi(
26     A: np.ndarray,
27     b: np.ndarray,
28     x0: np.ndarray,
29     k_max: int = 1000,
30     res_max: float = 1e-5,
31 ) -> Tuple[np.ndarray, List[float], List[float]]:
32     n = A.shape[0]
33     print("n:", n)
34     xk = np.copy(x0)
35     # "Iteration 0"
36     residuals = [residual(A, xk, b)]
37     errors = [err_func(xk)]
38     xkpl = np.zeros_like(xk)
39     for i in range(n):
40         for k in range(k_max):
41             for j in range(n):
42                 sigma = 0
43                 for j in range(n):
44                     if j != i:
45                         sigma += A[i, j] * xk[j]
46                 xkpl[i] = (b[i] - sigma) / A[i, i]
47             xk = xkpl
48             res = residual(A, xk, b)
49             err = err_func(xk)
50             residuals.append(res)
51             errors.append(err)
52
53             if res < res_max:
54                 print(f"Converged at iteration {k}")
55                 return xk, residuals, errors
56     else:
57         print(f"Failed to converge in {k_max} iterations!")
58         return xk, residuals, errors
59
60
61 def backward_Gauss_Seidel(
62     A: np.ndarray,
63     b: np.ndarray,
64     x0: np.ndarray,
65     k_max: int = 1000,
66     res_max: float = 1e-5,
67 ) -> Tuple[np.ndarray, List[float], List[float]]:
68     n = A.shape[0]
69     xk = x0
70     # "Iteration 0"
71     residuals = [residual(A, xk, b)]
72     errors = [err_func(xk)]
73     for k in range(k_max):
74         xkpl = xk
75         for i in reversed(range(n)):
76             # j > i
77             sum1 = np.sum(A[i, i + 1 :] * xk[i + 1 :])
78             # j < i
79             sum2 = np.sum(A[i, 0:i] * xk[:i])
80             xkpl[i] = (b[i] - sum1 - sum2) / A[i, i]
81
82         xk = xkpl
83
84         res = residual(A, xk, b)
85         err = err_func(xk)
86         residuals.append(res)
87         errors.append(err)
88
89         if res < res_max:

```

```

90         print(f"Converged at iteration {k}")
91         return xk, residuals, errors
92     else:
93         print(f"Failed to converged in {k_max} iterations!")
94         return xk, residuals, errors
95
96 def gen_PD_matrix(n: int) -> np.ndarray:
97     tmp = np.random.randn(n, n)
98     return tmp @ tmp.T
99
100
101 def main() -> None:
102     # Jacobi Plot
103     x_jac, jac_res, jac_err = Jacobi(A, b, np.array((1.0, 1.0, 1.0, 1.0)))
104     print(x_jac)
105     plt.figure(figsize=(8, 4.5))
106     plt.semilogy(jac_res, label="Residues")
107     plt.semilogy(jac_err, label="Error")
108     plt.xlabel("Iteration Number")
109     plt.ylabel("Error or Residue Values")
110     plt.title("Jacobi Method Convergence")
111     plt.legend()
112     plt.savefig("jacobi_iter.pdf", bbox_inches="tight")
113
114     # Backwards Gauss-Seidel
115     x_bgs, bgs_res, bgs_err = backward_Gauss_Seidel(
116         A, b, np.array((1.0, 1.0, 1.0, 1.0))
117     )
118     print(x_bgs)
119     plt.figure(figsize=(8, 4.5))
120     plt.semilogy(bgs_res, label="Residues")
121     plt.semilogy(bgs_err, label="Error")
122     plt.xlabel("Iteration Number")
123     plt.ylabel("Error or Residue Values")
124     plt.title("Backward Gauss-Seidel Method Convergence")
125     plt.legend()
126     plt.savefig("bgs_iter.pdf", bbox_inches="tight")
127
128     # SOR Investigation
129     omegas = np.linspace(0.01, 2.20, num=1000, endpoint=True)
130     plt.figure()
131     for i in range(20):
132         A_tmp = gen_PD_matrix(4)
133         D = np.diag(np.diag(A_tmp))
134         L = np.tril(A_tmp, -1)
135         radii = []
136         for omega in omegas:
137             M = (1.0 / omega) * D + L
138             radius = np.max(
139                 np.abs(np.linalg.eigvals(np.eye(4) - np.linalg.inv(M) @ A_tmp))
140             )
141             radii.append(radius)
142     plt.plot(omegas, radii)
143
144     plt.xlabel("Omega")
145     plt.ylabel("Spectral Radius")
146     plt.title(
147         f"Spectral Radius of the iteration matrix  $I - M^{-1}A$  for 20
148         random A matrices"
149     )
150     plt.savefig("sor-omegas.pdf", bbox_inches="tight")
151
152 if __name__ == "__main__":
153     main()

```

Bibliography

- [1] Stewart Venit. “The Convergence of Jacobi and Gauss-Seidel Iteration”. In: *Mathematics Magazine* 48.3 (May 1975), pp. 163–167.