



## Problem 1: Appreciate Cholesky Graphically

### Problem Statement



(a) Graph  $G_1$



(b) Graph  $G_2$

**a**

Find a matrix  $\mathbf{A} \in \mathbb{R}^{5 \times 5}$  s.t.  $\mathcal{G}(\mathbf{A}) = G_1$  in the figure above, and s.t.  $\mathbf{A} \succ 0$ . *Hint: Associate positive numbers or weights to each node and each edge of  $G_1$ . Define  $A_{ij} = A_{ji}$  as the negative of the weight on the edge  $(i, j)$  of  $G_1$ . Define  $A_{ii}$  to be the sum of the weight of node  $i$  and the weights of all edges connected to node  $i$ . Such a matrix is called a Laplacian matrix for  $G_1$ . Compute the eigenvalues to verify it is indeed PD.*

**b**

Compute the lower triangular Cholesky factor  $\mathbf{L} \in \mathbb{R}^{5 \times 5}$  with positive diagonals such that  $\mathbf{L}\mathbf{L}^T$  equals  $\mathbf{A}$  in part (a). Draw a graph on 5 nodes that describes the sparsity pattern of  $\mathbf{L}$ ; i.e. draw  $G(\mathbf{L})$ . Verify that  $G(\mathbf{L})$  is a chordal graph.

**c**

Notice that  $G_2$  in figure 1b is a chordal graph. Find a permutation matrix  $\mathbf{P}$  s.t. that when you compute the Cholesky factor  $\mathbf{L}'$  of  $\mathbf{P}\mathbf{A}\mathbf{P}^T$  with your  $\mathbf{A}$  from part (a), then  $G(\mathbf{L}') = G_2$ . *Hint: Eliminate the nodes of  $G_1$  in the sequence  $(5, 2, 3, 1, 4)$ . Encode that elimination order in a permutation matrix.*

**d(\*\*)**

Let  $\mathbf{X} \in \mathbb{R}^{n \times n}$  be a PD matrix and  $\mathbf{\Gamma} \in \mathbb{R}^{n \times n}$  be a lower triangular matrix with positive diagonals, s.t.  $\mathbf{X} = \mathbf{\Gamma}\mathbf{\Gamma}^T$ . Prove that  $G(\mathbf{X})$  is a subgraph of  $G(\mathbf{\Gamma})$  and  $G(\mathbf{\Gamma})$  is a chordal graph. *Hint: Use an induction argument on the size of the matrix. Characterize how the steps in Cholesky factorization shapes the sparsity pattern of  $\mathbf{\Gamma}$ . Finally, utilize the fact that if a node and a set of edges from that node is added to a chordal graph  $G$ , then the obtained graph is chordal IFF every two neighbors of the new node in  $G$  already shared an edge in  $G$ .*

**Solution****a****b****c****d**

## Problem 2: Split 'em matrices for Linear Systems

### Problem Statement

Consider a linear system of equations  $Ax = b$  where  $A \in \mathbb{R}^{n \times n}$  and  $\exists A^{-1}$ . The matrix splitting method splits  $A$  as  $A = M - N$  and successively solves the equation:

$$x^{k+1} = M^{-1} (Nx^k + b) \quad (1)$$

for  $k \geq 0$  to perform a fixed-point iteration in order to solve the linear system  $Ax = b$ . Identify  $L, D$  and  $U$  as the strictly lower triangular, diagonal, and strictly upper triangular part of  $A$ . Then,  $A = L + D + U$ . When  $M = L + D$ , then the fixed point iteration becomes:

$$x_i^{k+1} := \frac{1}{A_{ii}} \left[ b_i - \sum_{j < i} A_{ij} x_j^{k+1} - \sum_{j > i} A_{ij} x_j^k \right] \quad (2)$$

for  $i = 1, \dots, n$  and  $k \geq 0$ .

**a**

Write the update equation for each element of  $x$  as in equation 2 when  $M = D + U$  using backward induction to compute the iterates in equation 1 successively.

**b**

Call the method you devised in part (a) as the *backward* Gauss-Seidel method. Which among this method and Jacobi methods would you expect to converge faster to the solution of  $Ax = b$ ? Which one will you choose if distributed computation is involved, and why?

**c**

Implement both backward Gauss-Seidel and Jacobi methods to solve the linear system of equations  $Ax = b$  with

$$A := \begin{pmatrix} 10 & 5 & 3 & 4 \\ 4 & 10 & 2 & 1 \\ 1 & 3 & 8 & 2 \\ 1 & 6 & 3 & 9 \end{pmatrix}, \quad b := \begin{pmatrix} 4 \\ -5 \\ 4 \\ -11 \end{pmatrix}$$

starting with  $x^0 := (1, 1, 1, 1)^T$ . Plot the residues and errors in the successive iterates on a semilog plot. Iterate until the residue falls below  $10^{-5}$ . The error and residue, respectively, at iteration  $k$  are given by:

$$\epsilon^k := \|x - x^*\|, \quad r^k := \|b - Ax^k\|$$

where  $x^*$  is the solution to the linear system.

**d**

Recall that the method of successive over-relaxation (SOR) uses  $M = \frac{1}{\omega} D + L$ . Any matrix splitting method converges to the solution of a linear system *IF* the spectral radius of the matrix  $I - M^{-1}A$  is less than 1. Plot the spectral radius of  $I - M^{-1}A$  for SOR as a function of  $\omega$  in the range  $[0.01, 2.20]$  for 20 randomly generated PD  $A$  matrices on the same graph. Based on your plot, can you guess for what values of  $\omega$  SOR converges for PD matrices?

**Solution****a**

$$x_i^{k+1} := \frac{1}{A_{ii}} \left[ b_i - \sum_{j>i} A_{ij} x_j^{k+1} - \sum_{j<i} A_{ij} x_j^k \right]$$

**b**

It will depend. Gauss-Seidel methods use more “up to date” information, and so intuitively we would expect faster convergence. However, this is not a hard and fast rule, and there are counterexamples wherein one method converges and the other doesn’t, OR that one will converge faster in certain cases, or slower in others. See [1].

I think that using a typical scatter-gather scheme for distributed computation, Jacobi iteration would be preferred for distributed computing. Each element of  $x^{k+1}$  can be calculated in an embarrassingly parallel fashion (i.e. each element can be calculated in a fully independent fashion), whereas Gauss-Seidel will require iterations and sharing current estimates of  $x^{k+1}$  around, which will limit the ability to parallelize.

**c****d**

## Code

```

1  #!/usr/bin/env python3
2  import time
3  from typing import Callable, Iterable
4
5  import numpy as np
6  import torch
7
8  N_FACT = np.sqrt(2 * np.pi * 100) * (100 / np.e) ** 100
9
10
11 def naive_loop(x: Iterable) -> None:
12     sum_ = 0.0
13     for i in x:
14         sum_ += i
15
16
17 def numpy_vector_op(x: np.array) -> None:
18     sum_ = np.sum(x)
19
20
21 def torch_cpu_or_gpu(x: torch.Tensor) -> None:
22     sum_ = torch.sum(x)
23
24
25 def estimate_flops(
26     f: Callable[[int], float], n: int, lib: str = "numpy", iterations: int = 10
27 ) -> float:
28     times_ns = []
29     if lib == "native":
30         x = [1.0] * n
31     elif lib == "numpy":
32         x = np.ones(n, dtype=np.float64)
33     elif lib == "torch-cpu":
34         x = torch.ones(n, dtype=torch.float64)
35     elif lib == "torch-gpu":
36         x = torch.rand(n, dtype=torch.float64).to(torch.device("cuda:0"))
37     else:
38         raise RuntimeError(f"Invalid library {lib} requested!")
39
40     for _ in range(iterations):
41         torch.cuda.synchronize()
42         t0 = time.perf_counter_ns()
43         f(x)
44         t1 = time.perf_counter_ns()
45         times_ns.append(t1 - t0)
46
47     avg_time_ns = np.average(np.array(times_ns))
48
49     return n / (avg_time_ns / int(1e9))
50
51
52 def print_results(desc: str, flops: float) -> None:
53     naivest_loop_flops = estimate_flops(naive_loop, int(n), "native")
54     flop_time = 1.0 / flops
55     print(f"{desc}: {flops:.2E} FLOPS, or {flop_time:.2E} seconds for 1 FLOP")
56     print(f"Time for 100x100 Matrix: {(N_FACT/3.15576e+7)*flop_time} years")
57
58
59 if __name__ == "__main__":
60     n = 1e6
61     naivest_loop_flops = estimate_flops(naive_loop, int(n), "native")
62     print_results("Python list, naive for loop", naivest_loop_flops)
63
64     naive_loop_flops = estimate_flops(naive_loop, int(n), "native")
65     print_results("NumPy array, naive for loop", naive_loop_flops)
66
67     n = 1e8
68     numpy_flops = estimate_flops(numpy_vector_op, int(n), "numpy")
69     print_results("NumPy array, vectorized sum", numpy_flops)
70
71     torch_cpu_flops = estimate_flops(torch_cpu_or_gpu, int(n), "torch-cpu")
72     print_results("Torch Tensor (CPU), vectorized sum", torch_cpu_flops)
73
74     torch_gpu_flops = estimate_flops(torch_cpu_or_gpu, int(n), "torch-gpu")
75     print_results("Torch Tensor (GPU), vectorized sum", torch_gpu_flops)

```

# Bibliography

- [1] Stewart Venit. “The Convergence of Jacobi and Gauss-Seidel Iteration”. In: *Mathematics Magazine* 48.3 (May 1975), pp. 163–167.